Changepoint detection in an orchestra of musical instruments

Student 1118161

King's College London

September 11, 2018

Abstract

Music information retrieval tasks serve as faithful benchmarks for time-series analysis pipelines due to the availability of strongly labelled training data such as MusicNet. Q peak detection algorithm on a hand-crafted feature, hidden Markov models and causal convolutional neural networks are compared in their ability to detect changepoints in recordings of eleven orchestral instruments. Detections are evaluated quantitatively with precision-recall metrics.

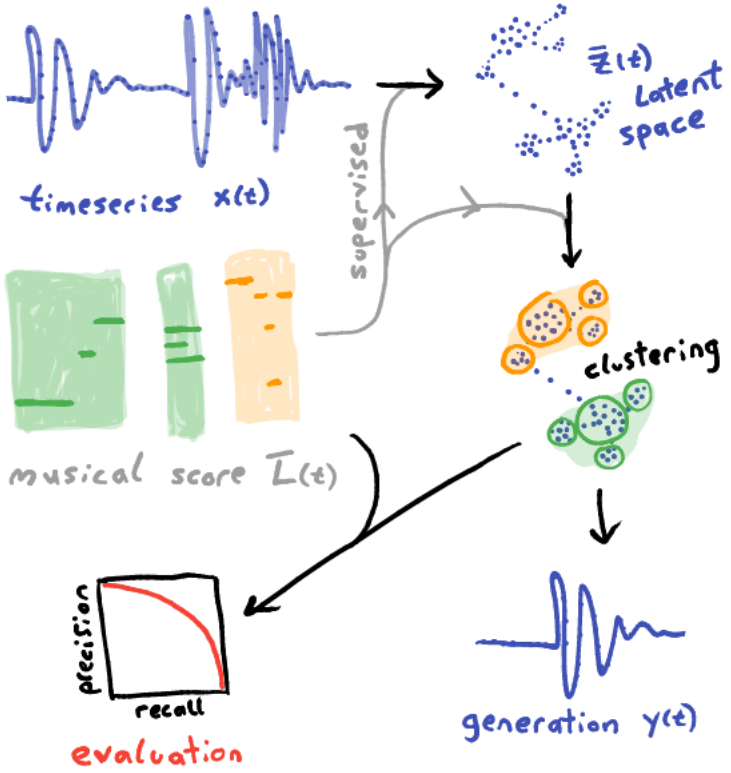

Figure 1: Summary of methodology showing all stages of the audio segmentation task. Each transition between sub-figures can be achieved with appropriate algorithms

1 Methodology Outline

1.1 Mapping time-series to latent space

The input data are single channel time-series points $\mathcal{D} = \{x(t_1) \dots x(t_N)\}$ sampled at frequency f from an underlying continuous state-time process $x(t)$, that is the oscillating sound waves emitted by a live orchestra.

For each time point there is a polyphonic score matrix $\mathbf{L}(t) \in \{0,1\}^{N \times L}$ which representing the N midi notes played by L instruments. Instrument activations $\bar{\mathrm{y}}(t) \in \{0,1\}^{L}$ and N note activations $\bar{\mathrm{n}}(t) \in \{0,1\}^{N}$ can be individually considered as well as only change-points between states. The most difficult task is to recover the polyphonic score matrix $\mathbf{L}(t) \in \{0,1\}^{N \times L}$ from the input signal $x(t)$. The score matrix is generally sparse as playing too many notes with too many instruments at the same time sounds terrible. Sparse labels

lead to sparse learning signals which hamper the convergence of machine learning algorithms. Thus we first look to recover changepoints and then the more densly labelled $\bar{\mathrm{n}}(t) \in \{0,1\}^N$ and $\bar{\mathrm{y}}(t) \in \{0,1\}^L$. From this we attempt to recover the full matrix in post-processing.

Within unsupervised methods this task is known as under-determined blind source separation. When the number of input channels equals to the number of sources this problem is fully determined and can be solved using independent component analysis [1] and other algorithms that search for sparse representations. Through the lens of supervised approaches this problem can be seen as an audio segmentation task. In recent years convolutional networks have demonstrated success in image segmentation [], whos architectures are adaptable to audio data.

The principal assumption in this task is that $x(t)$ is a linear superposition of sources and that in some latent space $\bar{\mathrm{z}}(t)$ these sources are linearly separable. The optimal latent space mapping f should minimise the cost function $J[f]$, which in general is some norm $\|\,\|$ of the error subject to some complexity measure $S[f]$. The data in the latent space should be separable by a suitable choice of unmixing matrix $\mathbf{W}$.

$$J[f] = \|\bar{\mathrm{y}}(t) - \mathbf{W}\,\bar{\mathrm{z}}(t)\| + S[f] \quad \text{where} \quad \bar{\mathrm{z}}(t) = f(x,t) \tag{1.1}$$

$$\exists\, \mathbf{W}, f \,:\, f = \underset{f'}{\operatorname{argmin}}\, J[f'] \tag{1.2}$$

Note that the map f is most likely not one-to-one respect to t. It is not a single time point but a time window that encodes which instrument is being played. While linear separability via $\mathbf{W}$ by is sufficient it is not necessary; one could instead simply look for clusters.

1.1.1 Wavelet transforms and spectral speed

The naive approach would be to attempt to guess the mapping f via suitable phenomenological agruments. It is reasonable to suppose that one could separate instruments based on their spectral signature at any moment in time. The decibel spectrogram $S(\omega, t)$ is obtained by taking the $\log_{10}$ absolute value of a short-time fourier transform.

$$S(\omega,t) := \log_{10}\left[\left|\int_{-\infty}^{\infty} x(\tau)\mathrm{e}^{-\mathrm{i}\omega\tau} h(\tau - t)\,\mathrm{d}\tau\,\right|\right] \tag{1.3}$$

A suitable window function $h(t)$ must be chosen; different choices lead to different amounts of spectral leakage and time-frequency domain resolution []. For demonstration purposes a 46ms wide Hann window is chosen with a stride of 3ms, which is on the order of the shortest

duration of a musical sound.

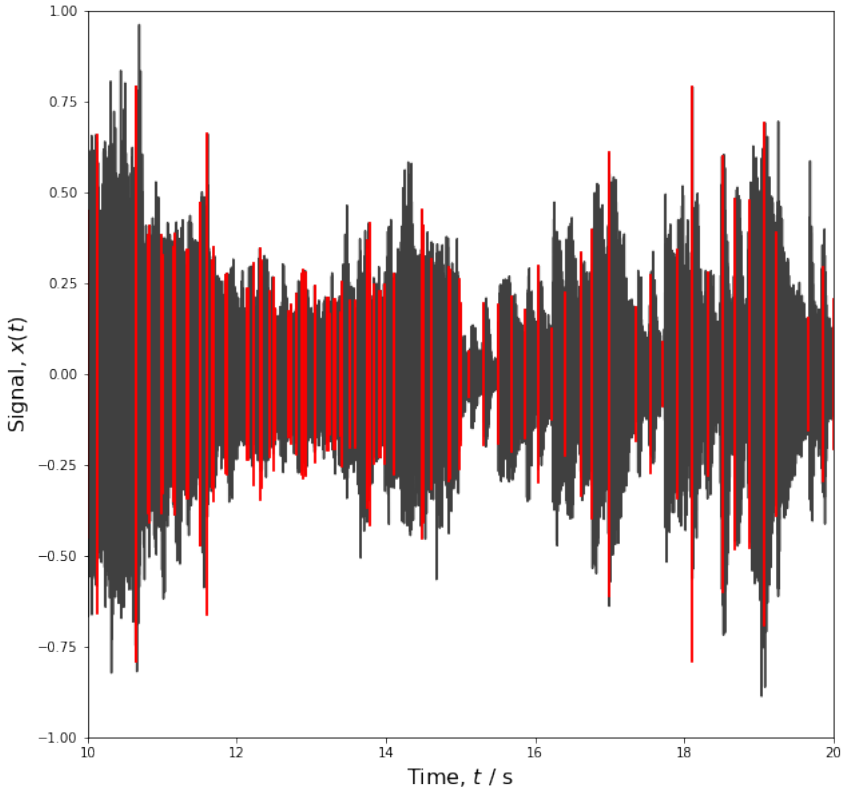

Figure 2: Input signal labelled with changepoints in polyphonic score $\mathbf{L}(t)$

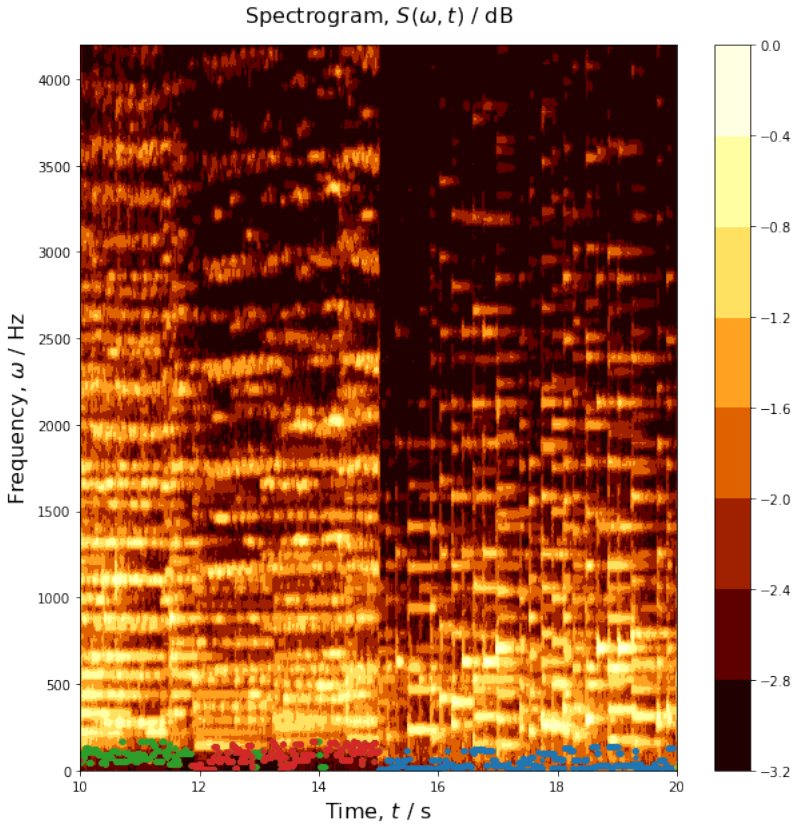

Figure 3: Spectrogram $S(\omega, t)$ with dominant frequency labelled by instrument $\bar{y}(t)$

The spectral signature of the input signal at a given time t can be represented as an F dimensional point, where F is the number of frequency bins up to a given cutoff frequency calculated in the spectrogram. The spectrogram in Figure 3 shows a promising variety of signatures within domain Ω that may differentiate between instruments. This observation motivates the calculation of spectral speed $v(t)$ defined as

$$v(t) := \frac{1}{|\Omega|} \int_{\Omega} \left| \partial_t S(\omega, t) \right| \mathrm{d}\omega \tag{1.4}$$

Figure 4 reveals that peaks in spectral speed $v(t)$ typically coincide with changepoints in notes and stepwise changes may indicate changepoints in instruments. Peaks can be detected using wavelet transform approaches [2] and true positives are counted when a detected peak lies within threshold time τ of a ground truth changepoint. Evaluations of 1000 clips of 5s from MusicNet in Figure 5 show that within 60ms a window approximately 60% of the changepoints can be recalled with 60% precision. The duration of a quaver amongst Beethoven symphonies range 40-125ms. This suggests that note changepoints more readily

distinguised in slower symphonies. Such a claim is supported by the example in 4, which shows a higher prevelance of false positives and negatives in the fast paced signal on the left-hand side. This approach is not precise enough to proceed with classification of the segments between changepoints.

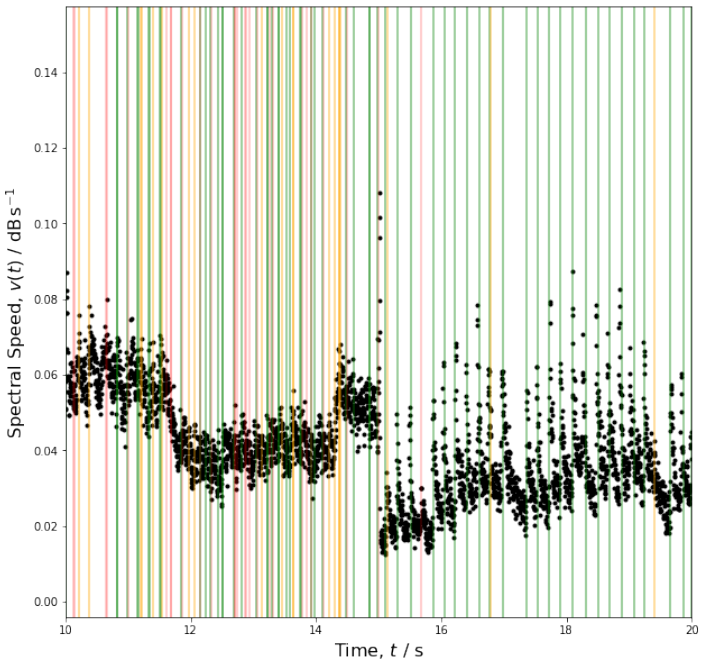

Figure 4: Spectral speed $v(t)$ labelled by true positive false positive and false negative changepoint detections

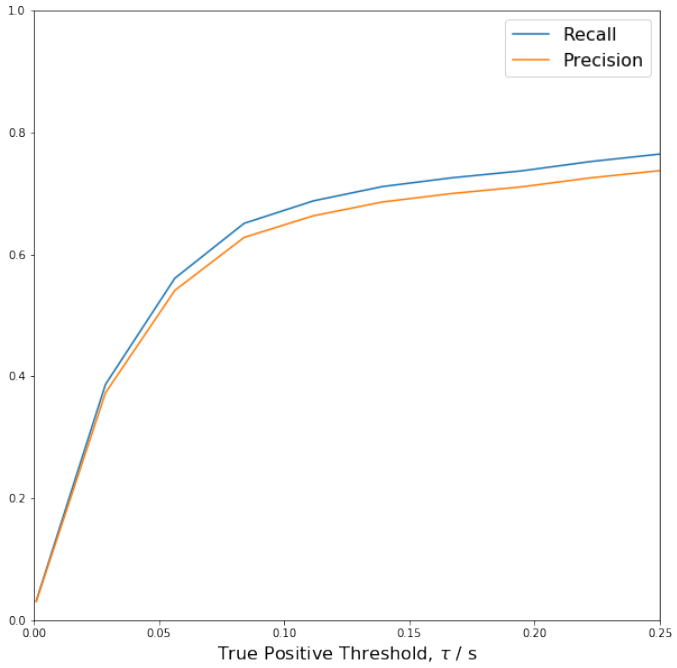

Figure 5: Changepoint precision-recall curves given by peaks in spectral speed $v(t)$

1.1.2 Markov models and expectation-maximisation

Hidden Markov models with Gaussian mixture emissions have been successful in the first generation of speech recognition piplines. Taking inspiration from these piplines, 11 separate bimodal Gaussian models are fitted via expectation-maximisation to 11 separate datasets. Each dataset contains 16 minute continuous concatinated spectrogram of a single instrument. Three hidden states maximise the average likelihood of across models. The classification of any new audio segment is given by the maximum likelihood model.

Figure 6 shows classification results for 1000 clips each lasting 100ms, whos ground truth labels are given by the dominant instrument in the clip. With an average total accuracy of 73% it outperforms the hand-crafted changepoint classifier in Section 1.1.1. Changepoints are detected as the boundary between any two classes, yielding the precision-recall curve in Figure 7

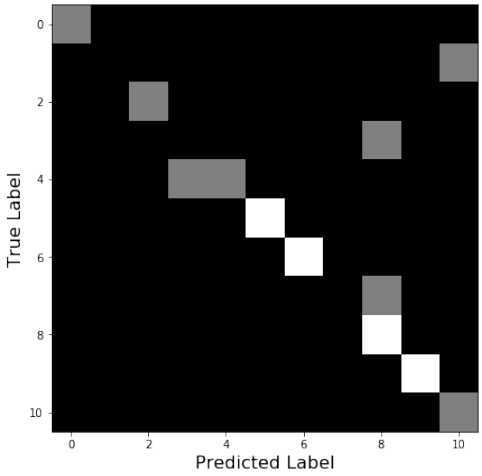

Figure 6: Instrument confusion matrix given by maximum likelihood Gaussian mixture Markov models for 100ms clip test set

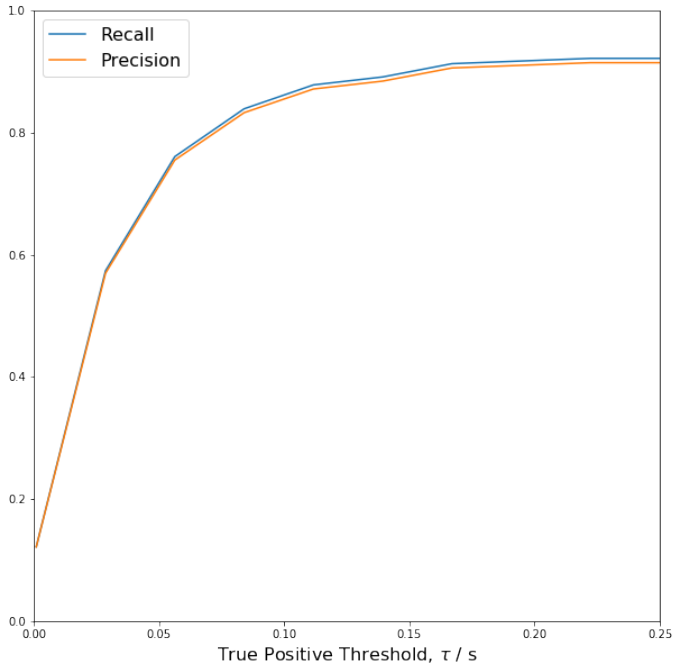

Figure 7: Changepoint precision-recall curves given by classification boundaries of maximum likelihood bimodal Gaussian mixture emissions Markov models

1.1.3 Feature extraction with causal convolutions

Convolutional architectures have become popular due to their ability to compress spatio-temporal information for discrimination and generation tasks [3, 4]. A causal convolutional network [5] — which encodes the arrow of time in its architecture — is trained for the audio segmentation task.

Often the performance of the classifiers is hampered by poor choices of feature space. Using the spectrogram as an input feature space in Sections 1.1.1 and 1.1.2 involves assumptions and arbitrary choices in the hyperparameters, such as the shape of the wavelet window $h(t)$. Figure 8 reveals a schematic of the architecture of WaveNet, where the inputs are the audio samples preceeding the next output sample. The arrow of time is encoded by the one-sided convolutions and the dilated/downsampled layers serve to reduce dimensionality and aggregate long range features in the waveform. In the original paper this architecture was trained to generate the next audio sample, based on previous samples. In our application we may ask for the instrument classification or probability of changepoint as an output.

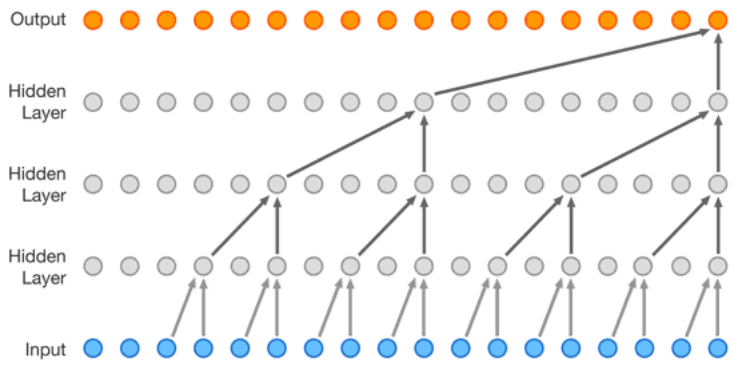

Figure 8: Causal convolutional architecture of WaveNet. Figure taken from [5]

Since this model takes the form of a computational graph, the back-propagation algorithm can be used to train it. This algorithm has an intrinsic advantage over expectation-maximisation as it can be optimally run on a GPU, which enables the training on much larger datasets. Large test and training sets are necessary for the models to generalise rather than overfit.

Due to time constraints and hardware limitations a smaller subset of MusicNet — 1 minute per instrument — was used as a training set to demonstrate overfitting. A standard adaptive momentum gradient descent method [6] was used and changepoints were detected as the boundary between any two classes on the same test set as in Section 1.1.2. Figure 9 reveals the convergence of the training protocol, giving rise to a very sensitive and nonspecific changepoint classifier as shown by precision-recall curves in Figure 10. The relaively poor precision could be accounted for by the noisy classification on the test set.

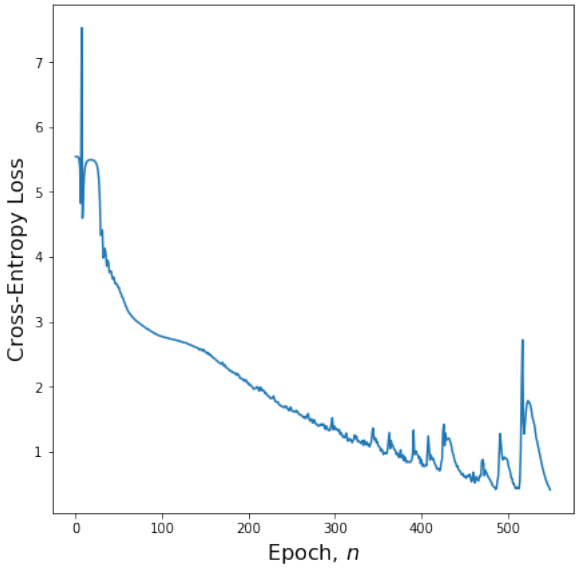

Figure 9: Training loss as a function of ADAM epoch, showing convergence of WaveNet

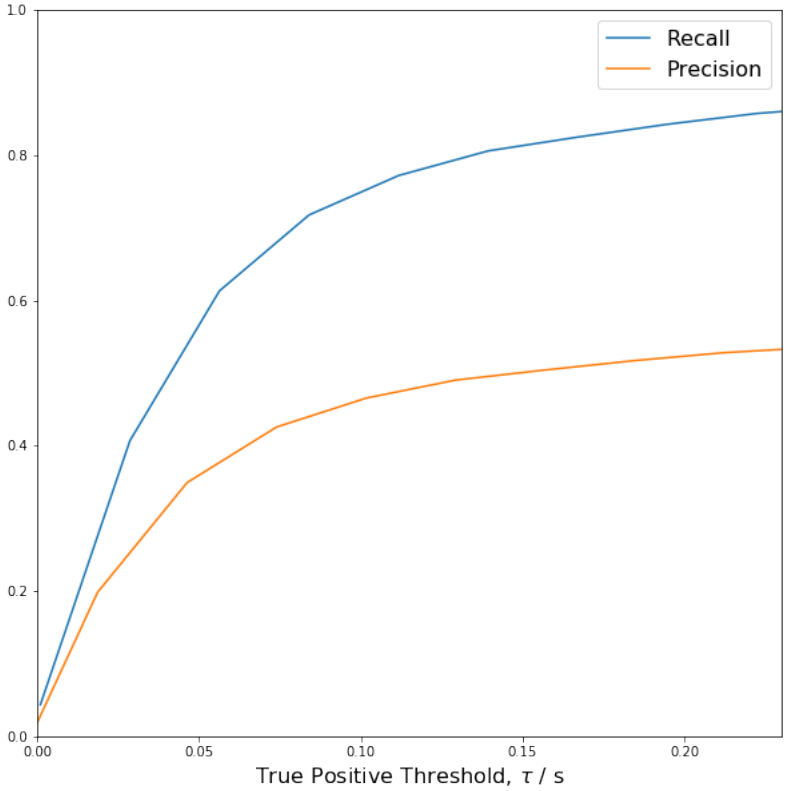

Figure 10: Changepoint precision-recall curves given by classification boundaries of WaveNet

References

[1] J. Platt and S. Haykin, "Information-Maximization Approach to Blind Separation and Blind Deconvolution," *Technology*, vol. 1159, no. 6, pp. 1129–1159, 1995.

[2] P. Du, W. A. Kibbe, and S. M. Lin, "Improved peak detection in mass spectrum by incorporating continuous wavelet transform-based pattern matching," *Bioinformatics*, vol. 22, pp. 2059–2065, sep 2006.

[3] A. v. d. Oord, N. Kalchbrenner, and K. Kavukcuoglu, "Pixel Recurrent Neural Networks," in *ICML, International Conference on Machine Learning*, vol. 48, pp. 1747–1756, 1 2016.

[4] I. Goodfellow, J. Pouget-Abadie, and M. Mirza, "Generative Adversarial Networks," *arXiv preprint*, pp. 1–9, 2014.

[5] A. van den Oord, S. Dieleman, H. Zen, K. Simonyan, O. Vinyals, A. Graves, N. Kalchbrenner, A. Senior, and K. Kavukcuoglu, "WaveNet: A Generative Model for Raw Audio," sep 2016.

[6] D. P. Kingma and J. Ba, "Adam: A Method for Stochastic Optimization," *arXiv*, dec 2014.